

Laboratorium 7: Restricted Boltzmann Machines i Deep Belief Networks

5 Ekstrakcja cech za pomocą Restricted Boltzmann Machine

- Wczytaj zbiór danych. Są to obrazy w skali szarości, przypisane do jednej z wielu klas, z pre-definiowanym podziałem na zbiór treningowy i testowy (w przypadku MNIST oraz Fashion-MNIST jest to 60000 rekordów treningowych i 10000 testowych).

```
import matplotlib.pyplot as plt
import tensorflow as tf
import numpy as np

(x_train, y_train), (x_test, y_test) =
tf.keras.datasets.mnist.load_data()
(f_x_train, f_y_train), (f_x_test, f_y_test) =
tf.keras.datasets.fashion_mnist.load_data()

k_x_train = np.load("kmnist-train-imgs.npz")['arr_0']
k_y_train = np.load("kmnist-train-labels.npz")['arr_0']
k_x_test = np.load("kmnist-test-imgs.npz")['arr_0']
k_y_test = np.load("kmnist-test-labels.npz")['arr_0']

# data for MNIST
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
print(f"Labels: {np.unique(y_train)}")

(60000, 28, 28)
(60000,)
(10000, 28, 28)
(10000,)
Labels: [0 1 2 3 4 5 6 7 8 9]

# data for FASHION-MNIST
print(f_x_train.shape)
print(f_y_train.shape)
print(f_x_test.shape)
print(f_y_test.shape)
print(f"Labels: {np.unique(f_y_train)}")

(60000, 28, 28)
(60000,)
(10000, 28, 28)
```

```

(10000,)
Labels: [0 1 2 3 4 5 6 7 8 9]

# data for Kuzushiji-MNIST
print(k_x_train.shape)
print(k_y_train.shape)
print(k_x_test.shape)
print(k_y_test.shape)
print(f"Labels: {np.unique(k_y_train)}")


(60000, 28, 28)
(60000,)
(10000, 28, 28)
(10000,)
Labels: [0 1 2 3 4 5 6 7 8 9]

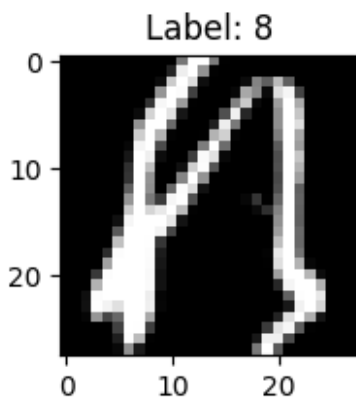
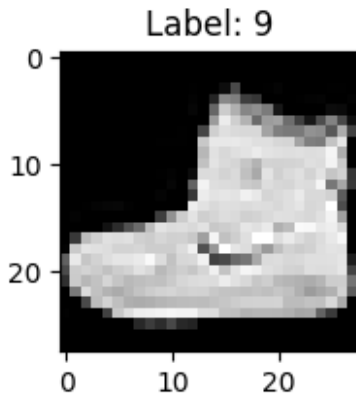
plt.figure(figsize=(2,2))
plt.imshow(x_train[0], cmap='gray')
plt.title(f"Label: {y_train[0]}")
plt.show()

plt.figure(figsize=(2,2))
plt.imshow(f_x_train[0], cmap='gray')
plt.title(f"Label: {f_y_train[0]}")
plt.show()

plt.figure(figsize=(2,2))
plt.imshow(k_x_train[0], cmap='gray')
plt.title(f"Label: {k_y_train[0]}")
plt.show()

```





- Dokonaj binaryzacji1 obrazów (zarówno zbiór treningowy, jak i testowy), a następnie wyświetl 10 przykładowych rekordów ze zbioru treningowego.
- Dokonaj konwersji danych na tablice numpy

```
from sklearn.preprocessing import binarize

def binarize_mnist(data):
    result = []
    for image in data:
        result.append(binarize(image, threshold=127))
    return result

def binarize_for_mnist():
    b_x_train = binarize_mnist(x_train)
    b_x_test = binarize_mnist(x_test)
    return np.array(b_x_train).astype(np.uint8),
    np.array(b_x_test).astype(np.uint8)

def binarize_for_f_mnist():
    b_f_x_train = binarize_mnist(f_x_train)
    b_f_x_test = binarize_mnist(f_x_test)
    return np.array(b_f_x_train).astype(np.uint8),
    np.array(b_f_x_test).astype(np.uint8)
```

```

def binarize_for_k_mnist():
    b_k_x_train = binarize_mnist(k_x_train)
    b_k_x_test = binarize_mnist(k_x_test)
    return np.array(b_k_x_train).astype(np.uint8),
    np.array(b_k_x_test).astype(np.uint8)

b_x_train, b_x_test = binarize_for_mnist()
b_f_x_train, b_f_x_test = binarize_for_f_mnist()
b_k_x_train, b_k_x_test = binarize_for_k_mnist()

y_train = np.array(y_train).astype(np.uint8)
y_test = np.array(y_test).astype(np.uint8)

f_y_train = np.array(f_y_train).astype(np.uint8)
f_y_test = np.array(f_y_test).astype(np.uint8)

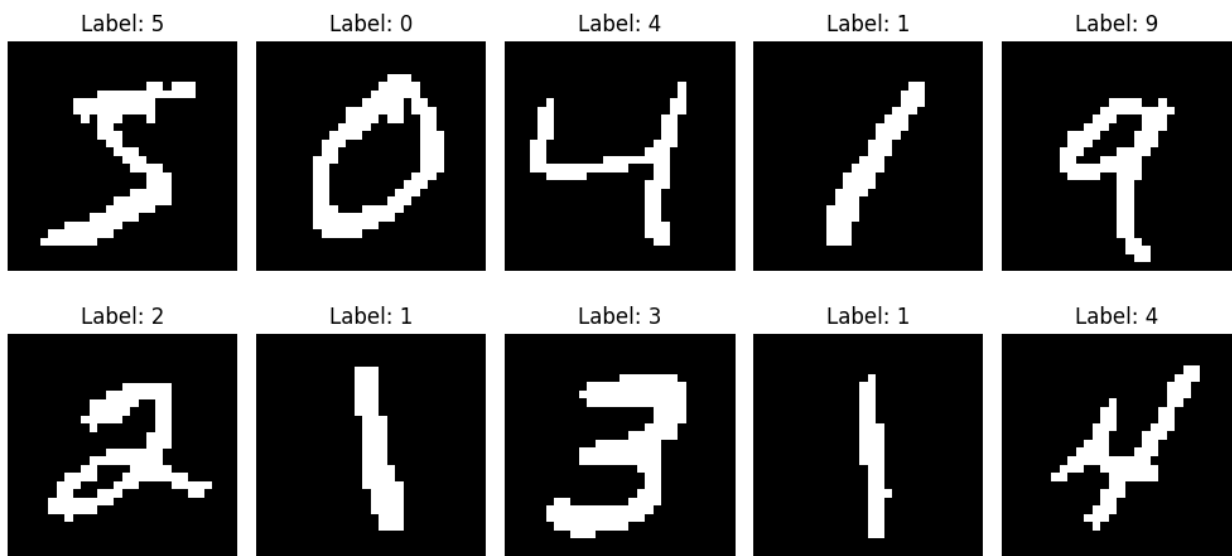
k_y_train = np.array(k_y_train).astype(np.uint8)
k_y_test = np.array(k_y_test).astype(np.uint8)

def plot_ten_images(x_data, y_data, lbl='Label'):
    plt.figure(figsize=(10, 5))
    for i in range(10):
        plt.subplot(2, 5, i + 1)
        plt.imshow(x_data[i], cmap='gray')
        plt.title(f"{lbl}: {y_data[i]}")
        plt.axis('off')


    plt.tight_layout()
    plt.show()

plot_ten_images(b_x_train, y_train)


```



```
plot_ten_images(b_f_x_train, f_y_train)
```



```
plot_ten_images(b_k_x_train, k_y_train)
```



```
# Prepare data
# Flatten images
b_x_train = b_x_train.reshape(-1, 28 * 28)
b_x_test = b_x_test.reshape(-1, 28 * 28)

b_f_x_train = b_f_x_train.reshape(-1, 28 * 28)
b_f_x_test = b_f_x_test.reshape(-1, 28 * 28)

b_k_x_train = b_k_x_train.reshape(-1, 28 * 28)
```

```

b_k_x_test = b_k_x_test.reshape(-1, 28 * 28)

# Scale images to [0, 1]
x_train = x_train.astype(np.float32)/255.0
x_train = x_train.reshape(-1, 28*28)

f_x_train = f_x_train.astype(np.float32)/255.0
f_x_train = f_x_train.reshape(-1, 28*28)

k_x_train = k_x_train.astype(np.float32)/255.0
k_x_train = k_x_train.reshape(-1, 28*28)

x_test = x_test.astype(np.float32) / 255.0
x_test = x_test.reshape(-1, 28*28)

f_x_test = f_x_test.astype(np.float32) / 255.0
f_x_test = f_x_test.reshape(-1, 28*28)

k_x_test = k_x_test.astype(np.float32) / 255.0
k_x_test = k_x_test.reshape(-1, 28*28)

```

- Zbuduj Pipeline złożony z: – BernoulliRBM do ekstrakcji cech – LogisticRegression jako klasyfikatora

```

from sklearn.neural_network import BernoulliRBM
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV

baseline_classifier = LogisticRegression()
pipe = Pipeline([('rbm', BernoulliRBM(n_components=50, n_iter=2)),
('regression', LogisticRegression())])
f_pipe = Pipeline([('rbm', BernoulliRBM(n_components=50, n_iter=2)),
('regression', LogisticRegression())])
k_pipe = Pipeline([('rbm', BernoulliRBM(n_components=50, n_iter=2)),
('regression', LogisticRegression())])

pipe.fit(b_x_train, y_train).score(b_x_test, y_test)

/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```

```
regression
    n_iter_i = _check_optimize_result(
```

0.8947

```
pipe2 = Pipeline([('rbm', BernoulliRBM(n_components=50, n_iter=10)),
('regression', LogisticRegression())])
pipe2.fit(b_x_train, y_train).score(b_x_test, y_test)
```

```
/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
    n_iter_i = _check_optimize_result(
```

0.8961

```
f_pipe.fit(b_f_x_train, f_y_train).score(b_f_x_test, f_y_test)
```

```
/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
    n_iter_i = _check_optimize_result(
```

0.7227

```
k_pipe.fit(b_f_x_train, f_y_train).score(b_f_x_test, f_y_test)
```

```
/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

0.7144

- Dostrój hiperparametry modelu korzystając z wyszukiwania siatkowego i walidacji krzyżowej (GridSearchCV). Siatkę parametrów możesz dobrać stosownie do zbioru danych i dostępnych zasobów obliczeniowych. Przykładowa siatka znajduje się poniżej.

```
param_grid = {
    'rbm__n_components': [70, 80, 90],
    'rbm__learning_rate': [0.05, 0.1],
    'rbm__batch_size': [10, 20],
    'rbm__n_iter': [5],
    'regression__C': [0.1, 0.5, 1.0]
}

grid_search = GridSearchCV(
    estimator=pipe,
    param_grid=param_grid,
    scoring='accuracy',
    cv=5,
    verbose=2,
    n_jobs=5
)

grid_search.fit(b_x_train, y_train)

print("Best params:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

print("Best params:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

Best params: {'rbm__batch_size': 20, 'rbm__learning_rate': 0.05,
'rbm__n_components': 90, 'rbm__n_iter': 5, 'regression__C': 1.0}
Best score: 0.9215833333333332

f_grid_search = GridSearchCV(
    estimator=f_pipe,
    param_grid=param_grid,
    scoring='accuracy',
    cv=5,
    verbose=2,
    n_jobs=5
)
```



```
f_grid_search.fit(b_f_x_train, f_y_train)

print("Best params:", f_grid_search.best_params_)
print("Best score:", f_grid_search.best_score_)

Best params: {'rbm__batch_size': 20, 'rbm__learning_rate': 0.05,
'rbm__n_components': 90, 'rbm__n_iter': 5, 'regression__C': 1.0}
Best score: 0.7785833333333334

k_grid_search = GridSearchCV(
    estimator=k_pipe,
    param_grid=param_grid,
    scoring='accuracy',
    cv=5,
    verbose=2,
    n_jobs=5
)

k_grid_search.fit(b_k_x_train, k_y_train)

print("Best params:", k_grid_search.best_params_)
print("Best score:", k_grid_search.best_score_)

Best params: {'rbm__batch_size': 10, 'rbm__learning_rate': 0.1,
'rbm__n_components': 90, 'rbm__n_iter': 5, 'regression__C': 0.5}
Best score: 0.8454666666666666
```

- Przedstaw i skomentuj otrzymane wartości hiperparametrów, a następnie wytrenuj zgodnie z nimi RBM i klasyfikator.

Komentarz do wyników hiperparametrów:

Dla zbiorów MNIST i F-MNIST batch-size zostały wybrany większy co daje większą stabilizację obliczeń, a learning-rate mniejszy więc uczenie będzie dokładniejsze. Ilość n_components dla każdego z przypadków została wybrana największa. Łatwiej dokonać ekstrakcji cech przy większej ilości komponentów. Wtedy uczenie staje się dokładniejsze. Ilość iteracji została ustalona na 5, ponieważ wtedy są szybsze obliczenia. Domyślna wartość to 10 i dla ostatecznych wyników taka została ustalona. Dla K-MNIST parametry batch-size, learning-rate i regression_C zostały wybrane nieco inne i dla nich best score wyszedł większy, ale po testowaniu tych parametrów to nie potwierdza się w poniższych wynikach (tutaj best score 0.845, a poniżej 0.7285), dlatego nie rozumiem jak to się stało.

```
pipe = Pipeline([('rbm',
BernoulliRBM(n_components=grid_search.best_params_['rbm__n_components'],
batch_size=grid_search.best_params_['rbm__batch_size'],
learning_rate=grid_search.best_params_['rbm__learning_rate'],
n_iter=10)), ('regression',
LogisticRegression(C=grid_search.best_params_['regression__C']))])
f_pipe = Pipeline([('rbm',
```

```

BernoulliRBM(n_components=f_grid_search.best_params_['rbm__n_component
s'], batch_size=f_grid_search.best_params_['rbm__batch_size'],
learning_rate=f_grid_search.best_params_['rbm__learning_rate'],
n_iter=10)), ('regression',
LogisticRegression(C=f_grid_search.best_params_['regression__C'])))
k_pipe = Pipeline([('rbm',
BernoulliRBM(n_components=k_grid_search.best_params_['rbm__n_component
s'], batch_size=k_grid_search.best_params_['rbm__batch_size'],
learning_rate=k_grid_search.best_params_['rbm__learning_rate'],
n_iter=10)), ('regression',
LogisticRegression(C=k_grid_search.best_params_['regression__C'])))
pipe.fit(b_x_train, y_train).score(b_x_test, y_test)

```

```

/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

0.9266

```
f_pipe.fit(b_f_x_train, f_y_train).score(b_f_x_test, f_y_test)
```

```

/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

0.779

```
k_pipe.fit(b_k_x_train, k_y_train).score(b_k_x_test, k_y_test)
```

0.7285

```
k_pipe_from_grid = k_grid_search.best_estimator_  
k_pipe_from_grid.score(b_k_x_test, k_y_test)
```

0.7238

- Dokonaj klasyfikacji na zbiorze testowym i przedstaw szczegółowo wyniki (accuracy, precision, recall, f1-score): – Na poziomie każdej klasy – Zbiorcze

```
from sklearn.metrics import accuracy_score, recall_score, f1_score,  
precision_score
```

```
def result_scores(y_true, y_pred):  
    res_acc = accuracy_score(y_true, y_pred, )  
    res_recall = recall_score(y_true, y_pred, average='micro')  
    res_f1 = f1_score(y_true, y_pred, average='micro')  
    res_prec = precision_score(y_true, y_pred, average='micro')
```

```
print(f"Results: ")  
print(f"\tAccuracy: {res_acc}")  
print(f"\tRecall: {res_recall}")  
print(f"\tF1-score: {res_f1}")  
print(f"\tPrecision: {res_prec}")
```

Extract classes

```
def extract_classes(y_data):  
    mnist_classes = np.unique(y_data)  
    indices_of_mnist_classes = {  
        "0": [],  
        "1": [],  
        "2": [],  
        "3": [],  
        "4": [],  
        "5": [],  
        "6": [],  
        "7": [],  
        "8": [],  
        "9": []  
    }  
  
    amount_of_elements = 0  
    for nr_class in mnist_classes:  
        indices_of_mnist_classes[str(nr_class)] = np.argwhere(y_data ==  
nr_class).reshape(-1)  
        amount_of_elements +=  
len(indices_of_mnist_classes[str(nr_class)])  
  
    assert amount_of_elements == len(y_data)  
  
    return indices_of_mnist_classes
```

```

y_test_classes = extract_classes(y_test)
f_y_test_classes = extract_classes(f_y_test)
k_y_test_classes = extract_classes(k_y_test)

def results_for_every_class(model, x_data, y_data):
    data_classes = np.unique(y_data)
    indices_classes = extract_classes(y_data)

    for single_class in data_classes:
        print(f"Class: {single_class}")
        y_pred = model.predict(x_data[indices_classes[str(single_class)]])
        result_scores(y_data[indices_classes[str(single_class)]], y_pred)

print(f"For MNIST dataset:")
print(f"Baseline: ")
baseline_classifier.fit(x_train, y_train)
y_pred_baseline = baseline_classifier.predict(x_test)
result_scores(y_test, y_pred_baseline)

print(f"For classes: ")
results_for_every_class(pipe, b_x_test, y_test)

print(f"For all:")
y_pred_all = pipe.predict(b_x_test)
result_scores(y_test, y_pred_all)

```

For MNIST dataset:
Baseline:

```

/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Results:

```

Accuracy: 0.9256
Recall: 0.9256
F1-score: 0.9256
Precision: 0.9256

```

For classes:

Class: 0

Results:

Accuracy: 0.9744897959183674
Recall: 0.9744897959183674
F1-score: 0.9744897959183674
Precision: 0.9744897959183674

Class: 1

Results:

Accuracy: 0.9823788546255506
Recall: 0.9823788546255506
F1-score: 0.9823788546255506
Precision: 0.9823788546255506

Class: 2

Results:

Accuracy: 0.9127906976744186
Recall: 0.9127906976744186
F1-score: 0.9127906976744186
Precision: 0.9127906976744186

Class: 3

Results:

Accuracy: 0.9089108910891089
Recall: 0.9089108910891089
F1-score: 0.9089108910891089
Precision: 0.9089108910891089

Class: 4

Results:

Accuracy: 0.9276985743380856
Recall: 0.9276985743380856
F1-score: 0.9276985743380856
Precision: 0.9276985743380856

Class: 5

Results:

Accuracy: 0.874439461883408
Recall: 0.874439461883408
F1-score: 0.874439461883408
Precision: 0.874439461883408

Class: 6

Results:

Accuracy: 0.9561586638830898
Recall: 0.9561586638830898
F1-score: 0.9561586638830898
Precision: 0.9561586638830898

Class: 7

Results:

Accuracy: 0.919260700389105
Recall: 0.919260700389105
F1-score: 0.919260700389105
Precision: 0.919260700389105

Class: 8

Results:

Accuracy: 0.893223819301848

```
Recall: 0.893223819301848
F1-score: 0.893223819301848
Precision: 0.893223819301848
```

Class: 9

Results:

```
Accuracy: 0.9058473736372646
Recall: 0.9058473736372646
F1-score: 0.9058473736372646
Precision: 0.9058473736372646
```

For all:

Results:

```
Accuracy: 0.9266
Recall: 0.9266
F1-score: 0.9266
Precision: 0.9266
```

```
print(f"For FASHION-MNIST dataset:")
print(f"Baseline: ")
baseline_classifier.fit(f_x_train, f_y_train)
f_y_pred_baseline = baseline_classifier.predict(f_x_test)
result_scores(f_y_test, f_y_pred_baseline)
```

```
print(f"For classes: ")
results_for_every_class(f_pipe, b_f_x_test, f_y_test)
```

```
print(f"For all:")
f_y_pred_all = f_pipe.predict(b_f_x_test)
result_scores(f_y_test, f_y_pred_all)
```

For FASHION-MNIST dataset:

Baseline:

```
/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs
failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Results:

```
Accuracy: 0.8438
Recall: 0.8438
F1-score: 0.8438
Precision: 0.8438
```

```
For classes:
Class: 0
Results:
  Accuracy: 0.744
  Recall: 0.744
  F1-score: 0.744
  Precision: 0.744
Class: 1
Results:
  Accuracy: 0.946
  Recall: 0.946
  F1-score: 0.946
  Precision: 0.946
Class: 2
Results:
  Accuracy: 0.623
  Recall: 0.623
  F1-score: 0.623
  Precision: 0.623
Class: 3
Results:
  Accuracy: 0.813
  Recall: 0.813
  F1-score: 0.813
  Precision: 0.813
Class: 4
Results:
  Accuracy: 0.738
  Recall: 0.738
  F1-score: 0.738
  Precision: 0.738
Class: 5
Results:
  Accuracy: 0.869
  Recall: 0.869
  F1-score: 0.869
  Precision: 0.869
Class: 6
Results:
  Accuracy: 0.4
  Recall: 0.4
  F1-score: 0.4
  Precision: 0.4
Class: 7
Results:
  Accuracy: 0.834
  Recall: 0.834
  F1-score: 0.834
  Precision: 0.834
```

Class: 8
Results:
 Accuracy: 0.914
 Recall: 0.914
 F1-score: 0.914
 Precision: 0.914

Class: 9
Results:
 Accuracy: 0.909
 Recall: 0.909
 F1-score: 0.909
 Precision: 0.909

For all:
Results:
 Accuracy: 0.779
 Recall: 0.779
 F1-score: 0.779
 Precision: 0.779

```
print(f"For K-MNIST dataset:")  
print(f"Baseline: ")  
baseline_classifier.fit(k_x_train, k_y_train)  
k_y_pred_baseline = baseline_classifier.predict(k_x_test)  
result_scores(k_y_test, k_y_pred_baseline)
```

```
print(f"For classes: ")  
results_for_every_class(k_pipe, b_k_x_test, k_y_test)
```

```
print(f"For all:")  
k_y_pred_all = k_pipe.predict(b_k_x_test)  
result_scores(k_y_test, k_y_pred_all)
```

For K-MNIST dataset:
Baseline:

```
/home/adrian/dev/um-labs/lab_RBM/.venv/lib/python3.11/site-packages/  
sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs  
failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression  
    n_iter_i = _check_optimize_result(
```

Results:
 Accuracy: 0.6965


```
Recall: 0.6965
F1-score: 0.6965
Precision: 0.6965
For classes:
Class: 0
Results:
  Accuracy: 0.785
  Recall: 0.785
  F1-score: 0.785
  Precision: 0.785
Class: 1
Results:
  Accuracy: 0.733
  Recall: 0.733
  F1-score: 0.733
  Precision: 0.733
Class: 2
Results:
  Accuracy: 0.642
  Recall: 0.642
  F1-score: 0.642
  Precision: 0.642
Class: 3
Results:
  Accuracy: 0.777
  Recall: 0.777
  F1-score: 0.777
  Precision: 0.777
Class: 4
Results:
  Accuracy: 0.694
  Recall: 0.694
  F1-score: 0.694
  Precision: 0.694
Class: 5
Results:
  Accuracy: 0.742
  Recall: 0.742
  F1-score: 0.742
  Precision: 0.742
Class: 6
Results:
  Accuracy: 0.762
  Recall: 0.762
  F1-score: 0.762
  Precision: 0.762
Class: 7
Results:
  Accuracy: 0.666
```

```
Recall: 0.666
F1-score: 0.666
Precision: 0.666
Class: 8
Results:
  Accuracy: 0.752
  Recall: 0.752
  F1-score: 0.752
  Precision: 0.752
Class: 9
Results:
  Accuracy: 0.732
  Recall: 0.732
  F1-score: 0.732
  Precision: 0.732
For all:
Results:
  Accuracy: 0.7285
  Recall: 0.7285
  F1-score: 0.7285
  Precision: 0.7285
```

- Czy nieliniowa ekstrakcja cech za pomocą RBM poprawia wyniki klasyfikacji w porównaniu z baseline (regresja logistyczna na surowych pikselach)?

Odpowiedź:

Dla zbiorów danych MNIST i K-MNIST wyniki klasyfikacji zostały poprawione. Dla F-MNIST nie udało się poprawić wyników. Dla poszczególnych cech było różnie, czasami były poprawy względem baseline-u, a czasami nie.

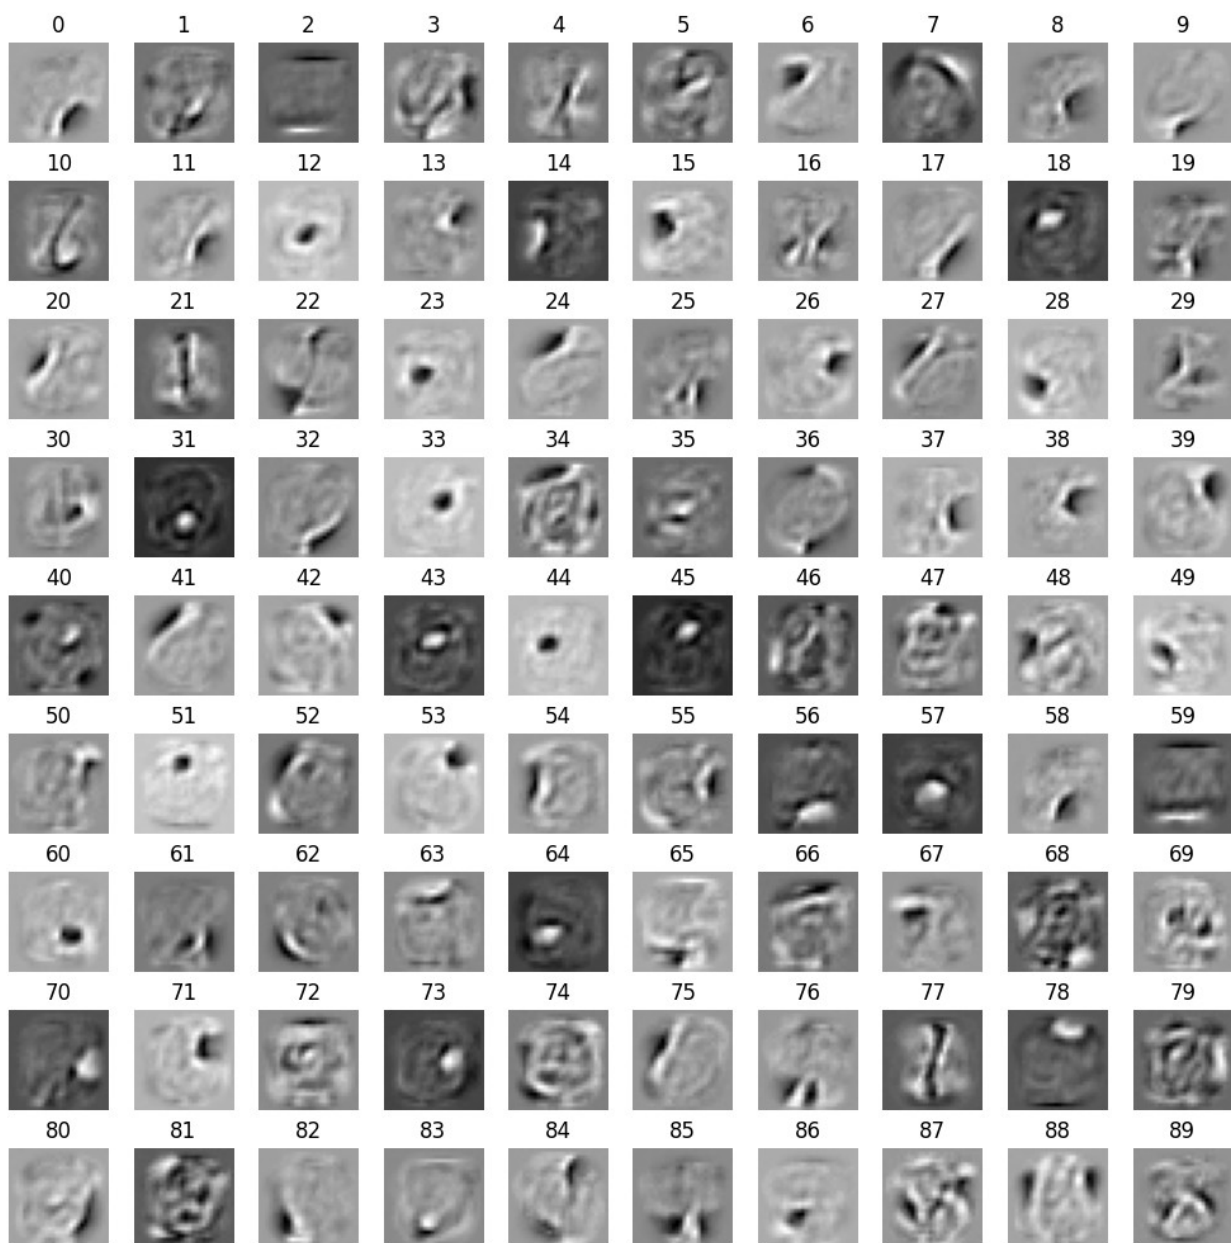
- Przedstaw wizualizację wszystkich ekstrahowanych cech ukrytych (n components obrazów odpowiadających wyuczonym wagom łączącym się z określonym elementem warstwy ukrytej). Spróbuj rozpoznać jakie wysokopoziomowe cechy obrazu są wzmacniane przez określony komponent.

```
def visualize_features(components):
    print(components.shape)
    image_shape = (28, 28)
    reshaped_components = components.reshape(-1, *image_shape)
    plt.figure(figsize=(10, 10))
    for i in range(reshaped_components.shape[0]):
        plt.subplot(reshaped_components.shape[0]//10, 10, i + 1)
        plt.imshow(reshaped_components[i], cmap='gray')
        plt.title(f"{i}")
        plt.axis('off')

    plt.tight_layout()
    plt.show()
```

Components for MNIST dataset:

```
visualize_features(pipe['rbm'].components_)  
(90, 784)
```




Components for F-MNIST dataset:

```
visualize_features(f_pipe['rbm'].components_)  
(90, 784)
```



Components for K-MNIST dataset:

```
visualize_features(k_pipe['rbm'].components_)  
(90, 784)
```



Dla poszczególnych komponentów widać czasami cechy jakie są wykrywane przez model np. są to ubrania jak koszule.

6 Hierarchiczna ekstrakcja cech za pomocą Deep Belief Network

1. Skonfiguruj 3 RBM w ten sposób, że warstwa ekstrakcji cech (hidden layer) RBM(i) staje się warstwą wejściową (visible layer) RBM(i + 1). RBMy trenowane są w sekwencyjnie: 1, 2, 3.

1. Rozmiar warstwy ukrytej jest pre-definiowany, przykładowo dla MNIST proponujemy kolejno: 256, 128, 64. Dla innych zbiorów danych rozmiary te możesz dopasować empirycznie, przy czym nie jest wymagane zastosowanie wyszukiwania siatkowego (duże wymagania obliczeniowe). Wartości pozostałych hiperparametrów mogą być podobne jak w poprzednim zadaniu.

```
def dbn_for_dataset(x_train_data, y_train_data, x_test_data,
y_test_data):
    rbm1 = BernoulliRBM(n_components=256, learning_rate=0.05, n_iter=7,
random_state=42)
    rbm2 = BernoulliRBM(n_components=128, learning_rate=0.05, n_iter=7,
random_state=42)
    rbm3 = BernoulliRBM(n_components=64, learning_rate=0.05, n_iter=7,
random_state=42)
    regression_classifier = LogisticRegression(max_iter=1000,
random_state=42)

    rbm1.fit(x_train_data)
    x_rbm1 = rbm1.transform(x_train_data)

    x_test_rbm1 = rbm1.transform(x_test_data)
    regression_classifier.fit(x_rbm1, y_train_data)
    y_pred_rbm1 = regression_classifier.predict(x_test_rbm1)

    rbm2.fit(x_rbm1)
    x_rbm2 = rbm2.transform(x_rbm1)

    x_test_rbm2 = rbm2.transform(x_test_rbm1)
    regression_classifier.fit(x_rbm2, y_train_data)
    y_pred_rbm2 = regression_classifier.predict(x_test_rbm2)

    rbm3.fit(x_rbm2)
    x_rbm3 = rbm3.transform(x_rbm2)

    x_test_rbm3 = rbm3.transform(x_test_rbm2)
    regression_classifier.fit(x_rbm3, y_train_data)
    y_pred_rbm3 = regression_classifier.predict(x_test_rbm3)

    print(f"Scores for first RBM: ")
    result_scores(y_test_data, y_pred_rbm1)
    print(f"Scores for first & second RBM: ")
    result_scores(y_test_data, y_pred_rbm2)
    print(f"Scores for DBN: ")
    result_scores(y_test_data, y_pred_rbm3)
```

1. Porównaj wyniki klasyfikacji uzyskane z zastosowaniem regresji logistycznej, dla poniższych wariantów ekstrakcji cech: • Baseline - regresja logistyczna na surowych pikselach • Pełna hierarchiczna ekstrakcja (wszystkie 3 RBMy) • Pierwszy RBM • Pierwszy i drugi RBM

```

def baseline(x_train_data, y_train_data, x_test_data, y_test_data):
    regression_classifier = LogisticRegression(max_iter=1000,
random_state=42)
    regression_classifier.fit(x_train_data, y_train_data)

    y_pred = regression_classifier.predict(x_test_data)
    print(f"Scores for baseline: ")
    result_scores(y_test_data, y_pred)

print(f"Baseline for MNIST:")
baseline(x_train, y_train, x_test, y_test)
print(f"Results for MNIST dataset: ")
dbn_for_dataset(b_x_train, y_train, b_x_test, y_test)

print(f"Baseline for F-MNIST:")
baseline(f_x_train, f_y_train, f_x_test, f_y_test)
print(f"Results for FASHION-MNIST dataset: ")
dbn_for_dataset(b_f_x_train, f_y_train, b_f_x_test, f_y_test)

print(f"Baseline for K-MNIST:")
baseline(k_x_train, k_y_train, k_x_test, k_y_test)
print(f"Results for K-MNIST dataset: ")
dbn_for_dataset(b_k_x_train, k_y_train, b_k_x_test, k_y_test)

```

Baseline for MNIST:

Scores for baseline:

Results:

Accuracy: 0.9263

Recall: 0.9263

F1-score: 0.9263

Precision: 0.9263

Results for MNIST dataset:

Scores for first RBM:

Results:

Accuracy: 0.9595

Recall: 0.9595

F1-score: 0.9595

Precision: 0.9595

Scores for first & second RBM:

Results:

Accuracy: 0.958

Recall: 0.958

F1-score: 0.958

Precision: 0.958

Scores for DBN:

Results:

Accuracy: 0.943

Recall: 0.943

F1-score: 0.943

Precision: 0.943

Baseline for F-MNIST:

Scores for baseline:

Results:

Accuracy: 0.8432

Recall: 0.8432

F1-score: 0.8432

Precision: 0.8432

Results for FASHION-MNIST dataset:

Scores for first RBM:

Results:

Accuracy: 0.802

Recall: 0.802

F1-score: 0.802

Precision: 0.802

Scores for first & second RBM:

Results:

Accuracy: 0.7647

Recall: 0.7647

F1-score: 0.7647

Precision: 0.7647

Scores for DBN:

Results:

Accuracy: 0.7232

Recall: 0.7232

F1-score: 0.7232

Precision: 0.7232

Baseline for K-MNIST:

Scores for baseline:

Results:

Accuracy: 0.6934

Recall: 0.6934

F1-score: 0.6934

Precision: 0.6934

Results for K-MNIST dataset:

Scores for first RBM:

Results:

Accuracy: 0.8073

Recall: 0.8073

F1-score: 0.8073

Precision: 0.8073

Scores for first & second RBM:

Results:

Accuracy: 0.7923

Recall: 0.7923

F1-score: 0.7923

Precision: 0.7923

Scores for DBN:

Results:

Accuracy: 0.7796

Recall: 0.7796
F1-score: 0.7796
Precision: 0.7796

1. Czy ustawianie ekstraktorów RBM jest korzystne dla każdego ze zbiorów danych? W jaki przypadku hierarchiczna ekstrakcja cech może być przydatna?

Ustawienie ekstraktorów jest korzystne dla zbiorów danych MNIST i K-MNIST. Nawet jeden RBM już znacząco poprawiał wyniki. Hierarchiczna ekstrakcja cech jest przydatna, gdy potrzebujemy badać cechy już w znalezionych cechach w danych.

7 ★ Restricted Boltzmann Machines: implementacja, testy rekonstrukcji

1. Zaimplementuj RBM wraz z algorytmem trenowania Contrastive Divergence. Możesz użyć następującego szablonu:

```
import time

def sigmoid(x):
    return 1.0 / (1.0 + np.exp(-x))

class RBM():
    def __init__(self, visible_dim, hidden_dim, learning_rate,
gifname='./images/rbm_training.gif'):
        self.visible_dim = visible_dim
        self.hidden_dim = hidden_dim
        self.learning_rate = learning_rate
        self.gif_name = gifname
        self.K = 1

        self.W = np.random.normal(0, 0.01, (visible_dim,
hidden_dim)).astype(np.float32)
        self.b = np.zeros(visible_dim)
        self.a = np.zeros(hidden_dim)

    def wake_phase(self, visible):
        preactivation = visible @ self.W
        preactivation = preactivation + np.broadcast_to(self.a,
preactivation.shape)
        activation = sigmoid(preactivation)
        binary_activation = activation >
np.random.rand(*activation.shape).astype(np.float32)
        return activation, binary_activation

    def dream_phase(self, hidden):
        preactivation = hidden @ self.W.T
        preactivation = preactivation + np.broadcast_to(self.b,
```

```

preactivation.shape)
    activation = sigmoid(preactivation)
    binary_activation = activation >
np.random.rand(*activation.shape).astype(np.float32)
    return activation, binary_activation

def train(self, batch):
    positive_visible = batch
    positive_hidden, negative_hidden = self.wake_phase(batch)
    negative_visible = self.dream_phase(negative_hidden)[1]

    for k in range(self.K):
        negative_hidden = self.wake_phase(negative_visible)[1]
        negative_visible = self.dream_phase(negative_hidden)[1]

    negative_hidden = self.wake_phase(negative_visible)[0]

    grad_positive = (positive_visible.T @ positive_hidden) /
batch.shape[0]
    grad_negative = (negative_visible.T @ negative_hidden) /
batch.shape[0]

    self.W += self.learning_rate * (grad_positive - grad_negative)
    return np.mean(np.abs(positive_visible[positive_visible >= 0] -
negative_visible[positive_visible >= 0]))

def fit(self, X, epochs=50, batch_dim=128, lr=0.1):
    self.learning_rate = lr
    test_set = X[0:2*batch_dim]
    for epoch in range(epochs):

        start = time.time()
        train_loss = 0
        s = 0
        for i in range(0, len(X) - batch_dim, batch_dim):
            batch = X[i:i+batch_dim]
            train_loss += self.train(batch)
            s+=1
        stop = time.time()

        reconstruction_error = self.reconstruction_error(test_set)
        print(f"Epoch {epoch+1}, time: {stop - start},
reconstruction_error: {reconstruction_error}, train_loss:
{train_loss/s}")
        self.plot_weights(epoch)

def reconstruction_error(self, X):
    error = X - self.reconstruct(X)
    return np.sum(np.square(error))/X.shape[0]

```

```

def reconstruct(self, X):
    preactivation = X @ self.W
    preactivation += np.broadcast_to(self.a, preactivation.shape)
    hidden = sigmoid(preactivation)
    hidden = hidden > np.random.rand(*hidden.shape).astype(np.float64)
    preactivation = hidden @ self.W.T
    preactivation += np.broadcast_to(self.b, preactivation.shape)
    visible = sigmoid(preactivation)
    return visible

def plot_weights(self, epoch):
    filters = np.reshape(np.transpose(self.W),
newshape=(self.hidden_dim, 28, 28))
    filters = np.clip(filters, -1, 1)
    fig = plt.figure(figsize=(5, 5))

    for i in range(self.hidden_dim):
        plt.subplot(8, self.hidden_dim//8, i + 1)
        plt.imshow(filters[i], cmap='gray')
        plt.axis('off')

    plt.tight_layout()
    plt.savefig(f'./images/image_at_epoch_{epoch:04d}.png')
    # plt.show()
    plt.close(fig)

```

1. Wytrenuj RBM na zbiorze Fashion-MNIST (część treningowa) dla warstwy ukrytej (hidden dim) o rozmiarze 40. Zastosuj domyślne wartości parametrów metody fit().

```

rbm_f_mnist = RBM(28*28, 40, 0.1)
rbm_f_mnist.fit(b_f_x_train, epochs=30, batch_dim=128)

/tmp/ipykernel_5746/441891556.py:4: RuntimeWarning: overflow
encountered in exp
    return 1.0 / (1.0 + np.exp(-x))

Epoch 1, time: 9.231838703155518, reconstruction_error:
77.62876892089844, train_loss: 37.321746341765895
Epoch 2, time: 9.52246880531311, reconstruction_error:
71.9843521118164, train_loss: 29.099820788942576
Epoch 3, time: 13.405994176864624, reconstruction_error:
69.47746276855469, train_loss: 27.572250060130212
Epoch 4, time: 12.423358917236328, reconstruction_error:
68.17606353759766, train_loss: 26.806279864960807
Epoch 5, time: 10.08058476448059, reconstruction_error:
66.76380157470703, train_loss: 26.355702223046677
Epoch 6, time: 9.121027708053589, reconstruction_error:
66.30818176269531, train_loss: 25.963610888816664
Epoch 7, time: 8.778710126876831, reconstruction_error:

```

66.05725860595703, train_loss: 25.653657565540264
Epoch 8, time: 10.14504861831665, reconstruction_error:
65.95828247070312, train_loss: 25.503535248233934
Epoch 9, time: 11.130028009414673, reconstruction_error:
65.7374267578125, train_loss: 25.414321438962197
Epoch 10, time: 6.259296655654907, reconstruction_error:
66.0661849975586, train_loss: 25.30433257412539
Epoch 11, time: 5.841235876083374, reconstruction_error:
65.78767395019531, train_loss: 25.25900343611003
Epoch 12, time: 5.721011161804199, reconstruction_error:
65.65464782714844, train_loss: 25.226255226900445
Epoch 13, time: 11.770118474960327, reconstruction_error:
65.73456573486328, train_loss: 25.19484576683237
Epoch 14, time: 14.629521608352661, reconstruction_error:
65.1994400024414, train_loss: 25.162717703768234
Epoch 15, time: 11.65986943244934, reconstruction_error:
65.46661376953125, train_loss: 25.15456944120796
Epoch 16, time: 9.923290491104126, reconstruction_error:
65.24038696289062, train_loss: 25.107378544616967
Epoch 17, time: 9.647526741027832, reconstruction_error:
65.67645263671875, train_loss: 25.112819478001544
Epoch 18, time: 11.533295631408691, reconstruction_error:
65.80133819580078, train_loss: 25.06313060079801
Epoch 19, time: 10.197265148162842, reconstruction_error:
66.16260528564453, train_loss: 25.10699702432055
Epoch 20, time: 11.775666952133179, reconstruction_error:
66.6102066040039, train_loss: 25.16769861114521
Epoch 21, time: 11.13505506515503, reconstruction_error:
66.48805236816406, train_loss: 25.23518558690504
Epoch 22, time: 9.345412969589233, reconstruction_error:
66.26042938232422, train_loss: 25.23296922184339
Epoch 23, time: 9.32714295387268, reconstruction_error:
66.40379333496094, train_loss: 25.275643836260485
Epoch 24, time: 11.3136146068573, reconstruction_error:
66.45555877685547, train_loss: 25.30430649072528
Epoch 25, time: 11.55119252204895, reconstruction_error:
66.61209869384766, train_loss: 25.34479380444087
Epoch 26, time: 22.16575002670288, reconstruction_error:
66.93470001220703, train_loss: 25.36642882438434
Epoch 27, time: 10.273731470108032, reconstruction_error:
66.44853973388672, train_loss: 25.364823492247453
Epoch 28, time: 6.184662342071533, reconstruction_error:
67.05076599121094, train_loss: 25.37330455770048
Epoch 29, time: 7.593412399291992, reconstruction_error:
66.68024444580078, train_loss: 25.34871210603632
Epoch 30, time: 6.177523612976074, reconstruction_error:
66.87451171875, train_loss: 25.346356125583245

1. Zwizualizuj jak zmieniają się wagi w kolejnych epokach treningu. W tym celu wywołuj metodę `plot_weights()` w metodzie `fit()`. Przedstaw wyniki w formie animowanego pliku GIF. Zapisz obserwacje dotyczące postępów treningu.

```
import imageio
import glob
import embed
def create_gif():
    with imageio.get_writer('images/rbm_training.gif', mode='I') as
writer:
    filenames = glob.glob('images/image*.png')
    filenames = sorted(filenames)
    for filename in filenames:
        image = imageio.imread(filename)
        writer.append_data(image)
        image = imageio.imread(filename)
        writer.append_data(image)
create_gif()
```

/tmp/ipykernel_7783/1815423369.py:9: DeprecationWarning: Starting with ImageIO v3 the behavior of this function will switch to that of `io.v3.imread`. To keep the current behavior (and make this warning disappear) use ``import imageio.v2 as imageio`` or call ``imageio.v2.imread`` directly.

image = imageio.imread(filename)

/tmp/ipykernel_7783/1815423369.py:11: DeprecationWarning: Starting with ImageIO v3 the behavior of this function will switch to that of `io.v3.imread`. To keep the current behavior (and make this warning disappear) use ``import imageio.v2 as imageio`` or call ``imageio.v2.imread`` directly.

image = imageio.imread(filename)

"RBM training"

Obserwacje:

Podczas treningu kształty na obrazkach są coraz smuklejsze i dokładniejsze, a cechy coraz bardziej się uwidaczniają.

Wnioski:

Podczas trenowania dochodzi do ekstrakcji cech i kształty są coraz lepiej rozpoznawane.

1. Wybierz 10 przykładów ze zbioru testowego (po jednym dla każdej klasy). Następnie, użyj RBM do rekonstrukcji obrazu testowego. Porównaj otrzymaną rekonstrukcję z obrazem oryginalnym: • Przedstaw wizualizację: obraz oryginalny vs. obraz odtworzony na podstawie wyuczonych wag RBM (dla każdej klasy) • Ilościowo: zastosuj Structural Similarity Index (SSIM).

```
from skimage.metrics import structural_similarity as ssim
```

```

f_y_test_classes = extract_classes(f_y_test)


def show_examples(different_classes, x_dataset, y_dataset):
    indices = [different_classes[str(i)][0] for i in range(10)]
    labels = y_dataset[indices]
    reconstructions = rbm_f_mnist.reconstruct(x_dataset[indices])
    reconstructions = np.reshape(reconstructions, newshape=(10, 28,
28))
    reconstructions = np.clip(reconstructions, -1, 1)
    original = np.reshape(x_dataset[indices], newshape=(10, 28, 28))
    print(f"Reconstructed images: ")
    plot_ten_images(reconstructions, labels)
    print(f"Original images: ")
    plot_ten_images(original, labels)
    ssim_indices = []
    differences = []
    for i in range(10):
        ssim_index, diff = ssim(reconstructions[i], original[i],
data_range=1, full=True)
        ssim_indices.append(round(ssim_index, 2))
        differences.append(diff)
    plot_ten_images(differences, ssim_indices, "SSIM_index")

show_examples(f_y_test_classes, b_f_x_test, f_y_test)


/tmp/ipykernel_5746/441891556.py:4: RuntimeWarning: overflow
encountered in exp
    return 1.0 / (1.0 + np.exp(-x))

```

Reconstructed images:



Original images:



1. Skomentuj uzyskane wyniki.

Komentarz:

Uzyskane wyniki przypominają kształty oryginalnych obrazków. Posiadają ich zarys, ale nie są w stanie odzwierciedlać dokładnych szczegółów. SSIM wykazuje podobieństwo na poziomie 0.4 - 0.5, więc są to podobne obrazy, ale tylko do pewnego stopnia.

8 Ekstrakcja cech za pomocą Autoencodera

1. Wczytaj zbiór danych. Dokonaj skalowania cech do przedziału [0, 1] oraz konwersji na tablice numpy (zarówno dane treningowe, jak i testowe).

```
from sklearn.model_selection import train_test_split
l = len(x_train)
ratio = 1
t_x_train, t_x_val, t_y_train, t_y_val =
train_test_split(x_train[:int(ratio*l)], y_train[:int(ratio*l)],
test_size=0.2, random_state=41)
t_f_x_train, t_f_x_val, t_f_y_train, t_f_y_val =
train_test_split(f_x_train, f_y_train, test_size=0.2, random_state=41)
t_k_x_train, t_k_x_val, t_k_y_train, t_k_y_val =
train_test_split(k_x_train, k_y_train, test_size=0.2, random_state=41)
```

1. Zbuduj i wytrenuj prosty Autoencoder. Przykładowa konfiguracja prostego Autoencodera: • Kodowanie: aktywacja ReLU • Dekodowanie: aktywacja sigmoidalna • Adam optimizer • Loss function: binary crossentropy • Encoding dimension: 196 (liczba ekstrahowanych cech) • Liczba epok treningu: > 30

```
class Autoencoder(tf.keras.Model):
    def __init__(self, encoding_dimension):
        super(Autoencoder, self).__init__()
        self.encoding_dimension = encoding_dimension

        self.encoder = tf.keras.Sequential(
            [
                # tf.keras.layers.Flatten(input_shape=input_shape),
                tf.keras.layers.Dense(784, activation='relu'),
                tf.keras.layers.Dense(128, activation='relu'),
                tf.keras.layers.Dense(64, activation='relu'),
                tf.keras.layers.Dense(32, activation='relu'),
                tf.keras.layers.Dense(self.encoding_dimension)
            ]
        )

        self.decoder = tf.keras.Sequential(
            [
                tf.keras.layers.Dense(self.encoding_dimension,
activation='relu'),
                tf.keras.layers.Dense(32, activation='relu'),
                tf.keras.layers.Dense(64, activation='relu'),
                tf.keras.layers.Dense(128, activation='relu'),
                tf.keras.layers.Dense(784, activation='sigmoid'),
                # tf.keras.layers.Reshape(target_shape=input_shape)
            ]
        )
```



```

def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

latent_dim = 196
batch_size = 32
epochs = 35

autoencoder = Autoencoder(latent_dim)
autoencoder.compile(optimizer='adam',
loss=tf.keras.losses.BinaryCrossentropy(from_logits=False),
metrics=[tf.keras.metrics.Accuracy])

2025-01-03 19:49:19.271215: E
external/local_xla/xla/stream_executor/cuda/cuda_driver.cc:152] failed
call to cuInit: INTERNAL: CUDA error: Failed call to cuInit: UNKNOWN
ERROR (303)

f_autoencoder = Autoencoder(latent_dim)
f_autoencoder.compile(optimizer='adam',
loss=tf.keras.losses.BinaryCrossentropy(from_logits=False),
metrics=[tf.keras.metrics.Accuracy])

k_autoencoder = Autoencoder(latent_dim)
k_autoencoder.compile(optimizer='adam',
loss=tf.keras.losses.BinaryCrossentropy(from_logits=False),
metrics=[tf.keras.metrics.Accuracy])

```

1. Przedstaw postępy trenowania wyświetlając wartości funkcji strat (loss) na zbiorach treningowych i testowych dla kolejnych epok.

```

autoencoder.fit(x_train, x_train,
                epochs=epochs,
                shuffle=True,
                batch_size=batch_size,
                validation_data=(x_test, x_test))

2025-01-03 19:50:10.130263: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:83]
Allocation of 188160000 exceeds 10% of free system memory.
2025-01-03 19:50:10.361451: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:83]
Allocation of 188160000 exceeds 10% of free system memory.

Epoch 1/35
1875/1875 _____ 38s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.2235 - val_accuracy: 0.0000e+00 - val_loss: 0.1367
Epoch 2/35
1875/1875 _____ 34s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.1349 - val_accuracy: 0.0000e+00 - val_loss: 0.1241

```

```
Epoch 3/35
1875/1875 _____ 37s 20ms/step - accuracy: 2.7851e-08 -
loss: 0.1223 - val_accuracy: 0.0000e+00 - val_loss: 0.1142
Epoch 4/35
1875/1875 _____ 35s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.1130 - val_accuracy: 0.0000e+00 - val_loss: 0.1096
Epoch 5/35
1875/1875 _____ 30s 16ms/step - accuracy: 7.0394e-08 -
loss: 0.1075 - val_accuracy: 2.5510e-07 - val_loss: 0.1048
Epoch 6/35
1875/1875 _____ 31s 17ms/step - accuracy: 8.7345e-09 -
loss: 0.1044 - val_accuracy: 0.0000e+00 - val_loss: 0.1025
Epoch 7/35
1875/1875 _____ 30s 16ms/step - accuracy: 0.0000e+00 -
loss: 0.1025 - val_accuracy: 1.2755e-07 - val_loss: 0.1008
Epoch 8/35
1875/1875 _____ 32s 17ms/step - accuracy: 1.4384e-07 -
loss: 0.1010 - val_accuracy: 3.8265e-07 - val_loss: 0.0993
Epoch 9/35
1875/1875 _____ 32s 17ms/step - accuracy: 1.1853e-07 -
loss: 0.0996 - val_accuracy: 1.2755e-07 - val_loss: 0.0986
Epoch 10/35
1875/1875 _____ 41s 17ms/step - accuracy: 1.3909e-07 -
loss: 0.0983 - val_accuracy: 8.9286e-07 - val_loss: 0.0972
Epoch 11/35
1875/1875 _____ 31s 17ms/step - accuracy: 3.6967e-07 -
loss: 0.0973 - val_accuracy: 3.8265e-07 - val_loss: 0.0964
Epoch 12/35
1875/1875 _____ 33s 17ms/step - accuracy: 1.1794e-06 -
loss: 0.0963 - val_accuracy: 2.4235e-06 - val_loss: 0.0960
Epoch 13/35
1875/1875 _____ 32s 17ms/step - accuracy: 1.3490e-06 -
loss: 0.0954 - val_accuracy: 1.5306e-06 - val_loss: 0.0949
Epoch 14/35
1875/1875 _____ 32s 17ms/step - accuracy: 3.6246e-06 -
loss: 0.0947 - val_accuracy: 1.6582e-06 - val_loss: 0.0945
Epoch 15/35
1875/1875 _____ 31s 16ms/step - accuracy: 4.1669e-06 -
loss: 0.0938 - val_accuracy: 4.3367e-06 - val_loss: 0.0942
Epoch 16/35
1875/1875 _____ 33s 17ms/step - accuracy: 5.9225e-06 -
loss: 0.0935 - val_accuracy: 5.8673e-06 - val_loss: 0.0930
Epoch 17/35
1875/1875 _____ 32s 17ms/step - accuracy: 8.1361e-06 -
loss: 0.0930 - val_accuracy: 4.9745e-06 - val_loss: 0.0927
Epoch 18/35
1875/1875 _____ 34s 18ms/step - accuracy: 9.9479e-06 -
loss: 0.0925 - val_accuracy: 3.6990e-06 - val_loss: 0.0926
Epoch 19/35
```

1875/1875 _____ 35s 19ms/step - accuracy: 1.3450e-05 -
loss: 0.0922 - val_accuracy: 2.3342e-05 - val_loss: 0.0922
Epoch 20/35
1875/1875 _____ 34s 18ms/step - accuracy: 1.7051e-05 -
loss: 0.0918 - val_accuracy: 1.8240e-05 - val_loss: 0.0935
Epoch 21/35
1875/1875 _____ 34s 18ms/step - accuracy: 2.0425e-05 -
loss: 0.0914 - val_accuracy: 5.9694e-05 - val_loss: 0.0929
Epoch 22/35
1875/1875 _____ 34s 18ms/step - accuracy: 2.7709e-05 -
loss: 0.0913 - val_accuracy: 2.7041e-05 - val_loss: 0.0913
Epoch 23/35
1875/1875 _____ 35s 18ms/step - accuracy: 3.3794e-05 -
loss: 0.0909 - val_accuracy: 3.0102e-05 - val_loss: 0.0915
Epoch 24/35
1875/1875 _____ 34s 18ms/step - accuracy: 3.9096e-05 -
loss: 0.0907 - val_accuracy: 4.2219e-05 - val_loss: 0.0913
Epoch 25/35
1875/1875 _____ 34s 18ms/step - accuracy: 4.3014e-05 -
loss: 0.0905 - val_accuracy: 5.6505e-05 - val_loss: 0.0911
Epoch 26/35
1875/1875 _____ 35s 18ms/step - accuracy: 5.3975e-05 -
loss: 0.0900 - val_accuracy: 5.9184e-05 - val_loss: 0.0908
Epoch 27/35
1875/1875 _____ 34s 18ms/step - accuracy: 5.8900e-05 -
loss: 0.0898 - val_accuracy: 6.0714e-05 - val_loss: 0.0901
Epoch 28/35
1875/1875 _____ 34s 18ms/step - accuracy: 6.6200e-05 -
loss: 0.0896 - val_accuracy: 5.7398e-05 - val_loss: 0.0894
Epoch 29/35
1875/1875 _____ 35s 19ms/step - accuracy: 6.5761e-05 -
loss: 0.0893 - val_accuracy: 5.5357e-05 - val_loss: 0.0901
Epoch 30/35
1875/1875 _____ 34s 18ms/step - accuracy: 7.7039e-05 -
loss: 0.0892 - val_accuracy: 5.3571e-05 - val_loss: 0.0895
Epoch 31/35
1875/1875 _____ 26s 14ms/step - accuracy: 8.1733e-05 -
loss: 0.0889 - val_accuracy: 8.7117e-05 - val_loss: 0.0906
Epoch 32/35
1875/1875 _____ 25s 13ms/step - accuracy: 8.1336e-05 -
loss: 0.0889 - val_accuracy: 4.8980e-05 - val_loss: 0.0891
Epoch 33/35
1875/1875 _____ 24s 13ms/step - accuracy: 9.4896e-05 -
loss: 0.0885 - val_accuracy: 6.3776e-05 - val_loss: 0.0890
Epoch 34/35
1875/1875 _____ 23s 12ms/step - accuracy: 1.0104e-04 -
loss: 0.0884 - val_accuracy: 8.2398e-05 - val_loss: 0.0889
Epoch 35/35

```
1875/1875 _____ 25s 13ms/step - accuracy: 1.1041e-04 -  
loss: 0.0882 - val_accuracy: 8.2781e-05 - val_loss: 0.0886
```

```
<keras.src.callbacks.history.History at 0x7f488af36910>
```

```
f_autoencoder.fit(f_x_train, f_x_train,  
                  epochs=epochs,  
                  shuffle=True,  
                  batch_size=batch_size,  
                  validation_data=(f_x_test, f_x_test))
```

```
2025-01-03 20:44:17.879751: W  
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:83]  
Allocation of 188160000 exceeds 10% of free system memory.
```

```
Epoch 1/35
```

```
1875/1875 _____ 34s 17ms/step - accuracy: 7.8733e-08 -  
loss: 0.3602 - val_accuracy: 1.2755e-07 - val_loss: 0.3042
```

```
Epoch 2/35
```

```
1875/1875 _____ 36s 19ms/step - accuracy: 3.7592e-06 -  
loss: 0.2998 - val_accuracy: 5.1020e-06 - val_loss: 0.2968
```

```
Epoch 3/35
```

```
1875/1875 _____ 38s 20ms/step - accuracy: 3.1669e-05 -  
loss: 0.2946 - val_accuracy: 3.2270e-05 - val_loss: 0.2940
```

```
Epoch 4/35
```

```
1875/1875 _____ 38s 20ms/step - accuracy: 5.6528e-05 -  
loss: 0.2909 - val_accuracy: 1.3265e-05 - val_loss: 0.2920
```

```
Epoch 5/35
```

```
1875/1875 _____ 37s 20ms/step - accuracy: 4.3741e-05 -  
loss: 0.2884 - val_accuracy: 3.3163e-06 - val_loss: 0.2889
```

```
Epoch 6/35
```

```
1875/1875 _____ 30s 16ms/step - accuracy: 1.2906e-05 -  
loss: 0.2863 - val_accuracy: 7.6531e-07 - val_loss: 0.2867
```

```
Epoch 7/35
```

```
1875/1875 _____ 30s 16ms/step - accuracy: 4.7483e-06 -  
loss: 0.2839 - val_accuracy: 1.7857e-06 - val_loss: 0.2854
```

```
Epoch 8/35
```

```
1875/1875 _____ 29s 16ms/step - accuracy: 3.4578e-06 -  
loss: 0.2823 - val_accuracy: 2.1684e-06 - val_loss: 0.2844
```

```
Epoch 9/35
```

```
1875/1875 _____ 39s 21ms/step - accuracy: 5.3652e-06 -  
loss: 0.2819 - val_accuracy: 3.4439e-06 - val_loss: 0.2835
```

```
Epoch 10/35
```

```
1875/1875 _____ 37s 20ms/step - accuracy: 2.8862e-06 -  
loss: 0.2809 - val_accuracy: 5.1020e-07 - val_loss: 0.2825
```

```
Epoch 11/35
```

```
1875/1875 _____ 30s 16ms/step - accuracy: 4.5398e-06 -  
loss: 0.2794 - val_accuracy: 2.5510e-07 - val_loss: 0.2819
```

```
Epoch 12/35
```

```
1875/1875 _____ 30s 16ms/step - accuracy: 5.2639e-06 -
```

```
loss: 0.2792 - val_accuracy: 3.4439e-06 - val_loss: 0.2815
Epoch 13/35
1875/1875 _____ 30s 16ms/step - accuracy: 3.9119e-06 -
loss: 0.2790 - val_accuracy: 1.2755e-06 - val_loss: 0.2810
Epoch 14/35
1875/1875 _____ 30s 16ms/step - accuracy: 4.6298e-06 -
loss: 0.2784 - val_accuracy: 2.5510e-07 - val_loss: 0.2810
Epoch 15/35
1875/1875 _____ 29s 16ms/step - accuracy: 4.3416e-06 -
loss: 0.2781 - val_accuracy: 8.9286e-07 - val_loss: 0.2799
Epoch 16/35
1875/1875 _____ 31s 16ms/step - accuracy: 6.2480e-06 -
loss: 0.2780 - val_accuracy: 1.1480e-06 - val_loss: 0.2804
Epoch 17/35
1875/1875 _____ 30s 16ms/step - accuracy: 6.3066e-06 -
loss: 0.2770 - val_accuracy: 1.9133e-06 - val_loss: 0.2800
Epoch 18/35
1875/1875 _____ 34s 18ms/step - accuracy: 7.7761e-06 -
loss: 0.2766 - val_accuracy: 2.5510e-06 - val_loss: 0.2789
Epoch 19/35
1875/1875 _____ 37s 20ms/step - accuracy: 1.0817e-05 -
loss: 0.2761 - val_accuracy: 2.0408e-06 - val_loss: 0.2788
Epoch 20/35
1875/1875 _____ 41s 22ms/step - accuracy: 8.5414e-06 -
loss: 0.2767 - val_accuracy: 4.8469e-06 - val_loss: 0.2785
Epoch 21/35
1875/1875 _____ 30s 16ms/step - accuracy: 7.7074e-06 -
loss: 0.2751 - val_accuracy: 1.5306e-06 - val_loss: 0.2786
Epoch 22/35
1875/1875 _____ 30s 16ms/step - accuracy: 1.4433e-05 -
loss: 0.2769 - val_accuracy: 2.6786e-06 - val_loss: 0.2786
Epoch 23/35
1875/1875 _____ 30s 16ms/step - accuracy: 1.5450e-05 -
loss: 0.2761 - val_accuracy: 8.4184e-06 - val_loss: 0.2776
Epoch 24/35
1875/1875 _____ 30s 16ms/step - accuracy: 2.3432e-05 -
loss: 0.2749 - val_accuracy: 1.9388e-05 - val_loss: 0.2782
Epoch 25/35
1875/1875 _____ 30s 16ms/step - accuracy: 2.7979e-05 -
loss: 0.2753 - val_accuracy: 2.8699e-05 - val_loss: 0.2774
Epoch 26/35
1875/1875 _____ 30s 16ms/step - accuracy: 3.6084e-05 -
loss: 0.2745 - val_accuracy: 2.0281e-05 - val_loss: 0.2776
Epoch 27/35
1875/1875 _____ 30s 16ms/step - accuracy: 4.1789e-05 -
loss: 0.2744 - val_accuracy: 3.7883e-05 - val_loss: 0.2776
Epoch 28/35
1875/1875 _____ 30s 16ms/step - accuracy: 3.5832e-05 -
loss: 0.2739 - val_accuracy: 1.9260e-05 - val_loss: 0.2774
```

```

Epoch 29/35
1875/1875 _____ 30s 16ms/step - accuracy: 4.2334e-05 -
loss: 0.2744 - val_accuracy: 1.2372e-05 - val_loss: 0.2770
Epoch 30/35
1875/1875 _____ 30s 16ms/step - accuracy: 3.9478e-05 -
loss: 0.2749 - val_accuracy: 2.0153e-05 - val_loss: 0.2770
Epoch 31/35
1875/1875 _____ 29s 16ms/step - accuracy: 5.1228e-05 -
loss: 0.2736 - val_accuracy: 1.5561e-05 - val_loss: 0.2768
Epoch 32/35
1875/1875 _____ 30s 16ms/step - accuracy: 5.3395e-05 -
loss: 0.2734 - val_accuracy: 2.0918e-05 - val_loss: 0.2769
Epoch 33/35
1875/1875 _____ 30s 16ms/step - accuracy: 5.2796e-05 -
loss: 0.2738 - val_accuracy: 2.3597e-05 - val_loss: 0.2770
Epoch 34/35
1875/1875 _____ 30s 16ms/step - accuracy: 4.2895e-05 -
loss: 0.2736 - val_accuracy: 1.4031e-05 - val_loss: 0.2764
Epoch 35/35
1875/1875 _____ 30s 16ms/step - accuracy: 5.4191e-05 -
loss: 0.2736 - val_accuracy: 2.0026e-05 - val_loss: 0.2765

<keras.src.callbacks.history.History at 0x7f48582c6910>

k_autoencoder.fit(k_x_train, k_x_train,
                  epochs=epochs,
                  shuffle=True,
                  batch_size=batch_size,
                  validation_data=(k_x_test, k_x_test))

Epoch 1/35
1875/1875 _____ 36s 19ms/step - accuracy: 0.0000e+00 -
loss: 0.2981 - val_accuracy: 0.0000e+00 - val_loss: 0.2867
Epoch 2/35
1875/1875 _____ 41s 19ms/step - accuracy: 0.0000e+00 -
loss: 0.2672 - val_accuracy: 0.0000e+00 - val_loss: 0.2696
Epoch 3/35
1875/1875 _____ 33s 18ms/step - accuracy: 6.8010e-09 -
loss: 0.2516 - val_accuracy: 0.0000e+00 - val_loss: 0.2600
Epoch 4/35
1875/1875 _____ 37s 20ms/step - accuracy: 0.0000e+00 -
loss: 0.2419 - val_accuracy: 0.0000e+00 - val_loss: 0.2515
Epoch 5/35
1875/1875 _____ 34s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.2342 - val_accuracy: 0.0000e+00 - val_loss: 0.2485
Epoch 6/35
1875/1875 _____ 34s 18ms/step - accuracy: 1.5823e-08 -
loss: 0.2308 - val_accuracy: 0.0000e+00 - val_loss: 0.2441
Epoch 7/35
1875/1875 _____ 33s 18ms/step - accuracy: 0.0000e+00 -

```

```
loss: 0.2265 - val_accuracy: 0.0000e+00 - val_loss: 0.2439
Epoch 8/35
1875/1875 _____ 34s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.2241 - val_accuracy: 0.0000e+00 - val_loss: 0.2402
Epoch 9/35
1875/1875 _____ 34s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.2222 - val_accuracy: 0.0000e+00 - val_loss: 0.2375
Epoch 10/35
1875/1875 _____ 34s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.2198 - val_accuracy: 0.0000e+00 - val_loss: 0.2356
Epoch 11/35
1875/1875 _____ 35s 19ms/step - accuracy: 6.3127e-08 -
loss: 0.2181 - val_accuracy: 1.2755e-07 - val_loss: 0.2345
Epoch 12/35
1875/1875 _____ 34s 18ms/step - accuracy: 1.2266e-08 -
loss: 0.2174 - val_accuracy: 0.0000e+00 - val_loss: 0.2343
Epoch 13/35
1875/1875 _____ 35s 19ms/step - accuracy: 0.0000e+00 -
loss: 0.2158 - val_accuracy: 0.0000e+00 - val_loss: 0.2321
Epoch 14/35
1875/1875 _____ 35s 19ms/step - accuracy: 0.0000e+00 -
loss: 0.2148 - val_accuracy: 0.0000e+00 - val_loss: 0.2307
Epoch 15/35
1875/1875 _____ 35s 19ms/step - accuracy: 3.6927e-09 -
loss: 0.2133 - val_accuracy: 0.0000e+00 - val_loss: 0.2293
Epoch 16/35
1875/1875 _____ 33s 18ms/step - accuracy: 0.0000e+00 -
loss: 0.2121 - val_accuracy: 0.0000e+00 - val_loss: 0.2291
Epoch 17/35
1875/1875 _____ 31s 17ms/step - accuracy: 1.0218e-07 -
loss: 0.2110 - val_accuracy: 0.0000e+00 - val_loss: 0.2283
Epoch 18/35
1875/1875 _____ 36s 19ms/step - accuracy: 2.9603e-08 -
loss: 0.2103 - val_accuracy: 0.0000e+00 - val_loss: 0.2274
Epoch 19/35
1875/1875 _____ 54s 29ms/step - accuracy: 0.0000e+00 -
loss: 0.2099 - val_accuracy: 0.0000e+00 - val_loss: 0.2285
Epoch 20/35
1875/1875 _____ 55s 29ms/step - accuracy: 5.4110e-08 -
loss: 0.2088 - val_accuracy: 1.2755e-07 - val_loss: 0.2270
Epoch 21/35
1875/1875 _____ 52s 28ms/step - accuracy: 0.0000e+00 -
loss: 0.2080 - val_accuracy: 0.0000e+00 - val_loss: 0.2275
Epoch 22/35
1875/1875 _____ 54s 29ms/step - accuracy: 1.0257e-07 -
loss: 0.2081 - val_accuracy: 0.0000e+00 - val_loss: 0.2262
Epoch 23/35
1875/1875 _____ 54s 29ms/step - accuracy: 3.3156e-08 -
loss: 0.2082 - val_accuracy: 0.0000e+00 - val_loss: 0.2263
```

```

Epoch 24/35
1875/1875 _____ 53s 28ms/step - accuracy: 1.3005e-08 -
loss: 0.2077 - val_accuracy: 0.0000e+00 - val_loss: 0.2258
Epoch 25/35
1875/1875 _____ 54s 29ms/step - accuracy: 0.0000e+00 -
loss: 0.2071 - val_accuracy: 0.0000e+00 - val_loss: 0.2256
Epoch 26/35
1875/1875 _____ 53s 28ms/step - accuracy: 5.4650e-08 -
loss: 0.2071 - val_accuracy: 0.0000e+00 - val_loss: 0.2251
Epoch 27/35
1875/1875 _____ 54s 29ms/step - accuracy: 4.1283e-08 -
loss: 0.2065 - val_accuracy: 0.0000e+00 - val_loss: 0.2253
Epoch 28/35
1875/1875 _____ 54s 29ms/step - accuracy: 0.0000e+00 -
loss: 0.2059 - val_accuracy: 0.0000e+00 - val_loss: 0.2246
Epoch 29/35
1875/1875 _____ 54s 29ms/step - accuracy: 0.0000e+00 -
loss: 0.2059 - val_accuracy: 0.0000e+00 - val_loss: 0.2248
Epoch 30/35
1875/1875 _____ 55s 29ms/step - accuracy: 3.7958e-09 -
loss: 0.2053 - val_accuracy: 0.0000e+00 - val_loss: 0.2261
Epoch 31/35
1875/1875 _____ 54s 29ms/step - accuracy: 0.0000e+00 -
loss: 0.2049 - val_accuracy: 0.0000e+00 - val_loss: 0.2243
Epoch 32/35
1875/1875 _____ 55s 29ms/step - accuracy: 0.0000e+00 -
loss: 0.2052 - val_accuracy: 0.0000e+00 - val_loss: 0.2238
Epoch 33/35
1875/1875 _____ 54s 29ms/step - accuracy: 6.2936e-08 -
loss: 0.2045 - val_accuracy: 0.0000e+00 - val_loss: 0.2241
Epoch 34/35
1875/1875 _____ 54s 29ms/step - accuracy: 1.3434e-07 -
loss: 0.2040 - val_accuracy: 0.0000e+00 - val_loss: 0.2234
Epoch 35/35
1875/1875 _____ 56s 30ms/step - accuracy: 1.0019e-07 -
loss: 0.2041 - val_accuracy: 0.0000e+00 - val_loss: 0.2234

<keras.src.callbacks.history.History at 0x7f48908f1210>

```

1. Hiperparametry Autoencodera możesz dobrać wykorzystując wyszukiwanie siatkowe (grid search).

```

from sklearn.model_selection import GridSearchCV
from scikeras.wrappers import KerasRegressor, KerasClassifier

param_grid = {
    'encoding_dimension': [32, 196],
    'batch_size': [32, 64, 128],
    'epochs': [30, 40, 50]
}

```



```

kerasClassifier = KerasClassifier(build_fn=autoencoder, verbose=0)
grid_search = GridSearchCV(estimator=kerasClassifier,
param_grid=param_grid, cv=3)

grid_search.fit(t_x_train, t_x_train)

print("Best Hyperparameters:", grid_search.best_params_)

```

Wyszukiwanie siatkowe nie zadziałało, dlatego opóściłem ten podpunkt.

1. Korzystając w wytrenowanego Autoencodera wygeneruj nowe cechy dla zbioru treningowego oraz zbioru testowego.

```

reconstructions_x_train = autoencoder.predict(x_train)
reconstructions_x_test = autoencoder.predict(x_test)

reconstructions_f_x_train = f_autoencoder.predict(f_x_train)
reconstructions_f_x_test = f_autoencoder.predict(f_x_test)

reconstructions_k_x_train = k_autoencoder.predict(k_x_train)
reconstructions_k_x_test = k_autoencoder.predict(k_x_test)

1875/1875 _____ 13s 7ms/step
313/313 _____ 2s 6ms/step
1875/1875 _____ 10s 5ms/step
313/313 _____ 2s 5ms/step
1875/1875 _____ 11s 6ms/step
313/313 _____ 1s 4ms/step

```

1. Wytrenuj klasyfikator LogisticRegression ustawiając solver='newton-cg'. Dokonaj predykcji na zbiorze testowym oraz porównaj uzyskane wyniki z wynikami z sekcji 5 i 6. Czy zmiana klasyfikatora na RandomForestClassifier poprawi dokładność klasyfikacji? Przedstaw swój komentarz do całościowych wyników klasyfikacji.

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

```

Dla mnist:

```

classifier = LogisticRegression(solver='newton-cg')
classifier.fit(x_train, y_train)

y_pred = classifier.predict(x_test)
y_pred_recon = classifier.predict(reconstructions_x_test)
print(f"accuracy_score baseline: {accuracy_score(y_pred, y_test)},
accuracy_score reconstructions: {accuracy_score(y_pred_recon,
y_test)}")

accuracy_score baseline: 0.926, accuracy_score reconstructions: 0.9308

```

Dla F-mnist:

```
f_classifier = LogisticRegression(solver='newton-cg')
f_classifier.fit(f_x_train, f_y_train)

y_pred = f_classifier.predict(f_x_test)
y_pred_recon = f_classifier.predict(reconstructions_f_x_test)
print(f"accuracy_score baseline: {accuracy_score(y_pred, f_y_test)},
accuracy_score reconstructions: {accuracy_score(y_pred_recon,
f_y_test)}")

accuracy_score baseline: 0.8434, accuracy_score reconstructions:
0.8336
```

Dla K-mnist:

```
k_classifier = LogisticRegression(solver='newton-cg')
k_classifier.fit(k_x_train, k_y_train)

y_pred = k_classifier.predict(k_x_test)
y_pred_recon = k_classifier.predict(reconstructions_k_x_test)
print(f"accuracy_score baseline: {accuracy_score(y_pred, k_y_test)},
accuracy_score reconstructions: {accuracy_score(y_pred_recon,
k_y_test)}")

accuracy_score baseline: 0.694, accuracy_score reconstructions: 0.7015

rf_classifier = RandomForestClassifier()
rf_classifier.fit(x_train, y_train)

y_pred_rf = rf_classifier.predict(x_test)
y_pred_rf_recon = rf_classifier.predict(reconstructions_x_test)
print(f"accuracy_score baseline: {accuracy_score(y_pred_rf, y_test)},
accuracy_score reconstructions: {accuracy_score(y_pred_rf_recon,
y_test)}")

accuracy_score baseline: 0.969, accuracy_score reconstructions: 0.9285

f_rf_classifier = RandomForestClassifier()
f_rf_classifier.fit(f_x_train, f_y_train)
y_pred_rf = f_rf_classifier.predict(f_x_test)
y_pred_rf_recon = f_rf_classifier.predict(reconstructions_f_x_test)
print(f"accuracy_score baseline: {accuracy_score(y_pred_rf,
f_y_test)}, accuracy_score reconstructions:
{accuracy_score(y_pred_rf_recon, f_y_test)}")

accuracy_score baseline: 0.878, accuracy_score reconstructions: 0.82

k_rf_classifier = RandomForestClassifier()
k_rf_classifier.fit(k_x_train, k_y_train)
y_pred_rf = k_rf_classifier.predict(k_x_test)
```

```
y_pred_rf_recon = k_rf_classifier.predict(reconstructions_k_x_test)
print(f"accuracy_score baseline: {accuracy_score(y_pred_rf,
k_y_test)}, accuracy_score reconstructions:
{accuracy_score(y_pred_rf_recon, k_y_test)}")
```

```
accuracy_score baseline: 0.8553, accuracy_score reconstructions:
0.7156
```

Wnioski:

Klasyfikator random forest dla baseline-u zwracał lepsze wyniki niż logistic regression. Ale dla rekonstrukcji logistic regression ma lepsze wyniki. Porównując z podpunktem 5 to MNIST i F-MNIST miały lepszy wynik, a K-MNIST osiągnął słabszy. Porównując z podpunktem 6 to MNIST i K-MNIST miały słabszy wynik, a F-MNIST osiągnął lepszy. RBM i autoencoder poprawiły wyniki klasyfikacji. Ekstrakcja cech pomaga w wyszukiwaniu wzorców i klasyfikacji danych.

1. * Zaproponuj własną architekturę głębokiego Autoencodera wykorzystującego filtry konwolucyjne. Nowe podejście do ekstrakcji cech powinno poprawić dokładność klasyfikacji na wszystkich zbiorach danych.

Ten podpunkt został zrealizowany w oddzielnym notebooku o nazwie:
task7_Adrian_Chrobot.ipynb